

EVALUATING CLAIMS, NOT JUST SYSTEMS

A one-cycle audit of what the evidence does—and does not—support

The project contributes a traceable claim-level map, not a new ARC solver, benchmark score, or validated universal methodology. Its central lesson: mechanism occurrence, local utility, comparator superiority, and external generalization are separate questions.

THE CLAIM LADDER

An organizing lens for separating evidence questions — not a validated measurement scale

01 MECHANISM PRESENT

Did the proposed object or operation occur?

Example ceiling: reflection / rotation skill objects appeared in controlled synthetic tasks.

02 TASK-RELEVANT UTILITY

Did it contribute to a relevant outcome in that setting?

Some reuse and simulator-step savings observed; local utility does not prove policy value.

03 INCREMENTAL SUPERIORITY

Did it beat appropriate, sufficiently strong comparators?

Transfer-aware ranking tied simpler controls; skill-creation policies did not clear superiority gates.

04 EXTERNAL GENERALIZATION

Did it hold beyond the development setting?

The ARC-AGI-2 adapter had very low fit coverage and no accepted downstream objects.

NO OVERALL SCORE Each claim can stop at a different evidence ceiling.

THE DISPOSITION MAP

Each conclusion is reported at its own evidence ceiling; dispositions should not be collapsed into one success/failure grade.



THE RESEARCH-CYCLE MAP

Read each row across the six stages. Synthetic evidence, capability metrics, bounded ARC adaptation, and the BL1 prior-art stop are kept distinct.

Research Cycle 1: from claim to scientific disposition

Read across each row: hypothesis → mechanism → utility → comparator → generalization → disposition.

	1 · HYPOTHESIS	2 · MECHANISM	3 · UTILITY	4 · COMPARATOR	5 · GENERALIZATION	6 · DISPOSITION
ASFT patterns / skills	Recurring structure can support reusable skills and transfer.	Reflection/rotation patterns and skill objects observed synthetically.	Some reuse and simulator step savings observed in the tested generators.	Ranking tied frequency / no-transfer; creation had no superiority gate pass.	ARC adapter: very low fit coverage; no accepted downstream objects.	**PARTIAL / BOUNDED** Synthetic mechanism and local utility; superiority and ARC transfer unproven.
Capability growth metrics	Metric growth indicates a growing capability portfolio.	Positive CGR/IG/IGR in one Phase 13E synthetic cohort.	Mixed metric path: negative IE trend; composition and utility caveats.	Cross-phase CGR and baseline contrasts mixed; not a stable increment.	No validated capability group on the ARC adapter.	**CONSTRUCT UNVALIDATED** One positive cohort does not establish robust capability or intelligence growth.
BL1 boundary learning	Counterexamples refine applicability boundaries while preserving coverage.	Not implemented; no contexts generated.	Not measured.	No empirical comparison.	Not tested.	**P0 PRIOR-ART STOP** Program terminated before experiment; no efficacy conclusion.

Submission-scope note: T2T / S3 is excluded from this version because its archived completion timestamp cannot be reconciled with the manuscript date.
The immutable archive is unchanged; this figure does not report the S3 score or disposition.

SCOPE & PROVENANCE

WHAT IS IN SCOPE — AND WHAT IS NOT

A deliberately bounded retrospective case study; immutable archive preserved

INCLUDED IN THIS WRITE-UP

- **ASFT**
  Saved Phase 13–14 evidence synthesis
- **BL1**
  Bounded P0 prior-art decision; no experiment
- **Cycle closure**
  Claim dispositions and research-cycle decision

EXCLUDED FROM EVIDENCE

T2T / S3

An unresolved timestamp discrepancy prevents use in this submission version.

The immutable archive remains unchanged. No S3 result or disposition is used.

NO NEW EXPERIMENTS • NO POST-HOC RESCUE

Important exclusion: T2T/S3 is excluded because its immutable execution timestamp cannot be reconciled with the manuscript date. The archive is unchanged; no S3 result or disposition is used. The literature screen is purposive, not systematic; the case supports no field-wide prevalence claim.

WHAT THE CASE SUPPORTS

Claim	Disposition	Boundary
Pattern discovery	Partial support	Synthetic generators; no external validation
Skill formation / reuse	Mixed	Some local utility; superiority not established
Capability growth	Unvalidated construct	Mixed phases; no validated ARC capability group
ARC transfer	Not demonstrated	Low-coverage adapter; non-activation, not falsification
BL1 boundary learning	Terminated at P0	Prior-art stop before implementation or experiment

READ THE FULL PAPER

The attached manuscript documents scope, source records, the claim ladder, comparative results, limitations, citations, and the exact provenance boundary. It reports no new experimental data or post-hoc rescue. The accompanying visual assets are explanatory summaries of that manuscript.

Title: Evaluating Claims, Not Just Systems — A Longitudinal Audit of an ARC Research Program